

Week1_Assignment

August 2, 2019

1 Exercise 1 of Week 1 Assignement

Consider the text available in this website <http://shakespeare.mit.edu/allswell/full.html>- to answer the following questions - (a) Find the frequency the word “BERTRAM” (all in caps, ignore font types such as bold, italics, etc.) that appears at the start of the sentence? - (b) Is the ratio of the frequency of

$$\frac{BERTRAN}{Bertram} > 0$$

- (c) What is the average sentence length of this document? - (d) What is the vocabulary (unique set of words) size? Ignore alpha- numeric and numeric terms, if available - (e) Find the total count of the words that end with the exclamation mark(!)

```
In [1]: import re
import urllib.request
from bs4 import BeautifulSoup
import nltk
from nltk.tokenize import word_tokenize # for splitting the text into words
from nltk.tokenize import sent_tokenize # for splitting the text into sentences
from nltk.tokenize import RegexpTokenizer # for extracting text ending with !
from nltk.corpus import inaugural # for average sentence length
```

1.1 Download the Text

```
In [2]: url = 'http://shakespeare.mit.edu/allswell/full.html'
html = urllib.request.urlopen(url).read().decode('utf-8')
soup = BeautifulSoup(html)
raw = BeautifulSoup(html, 'html.parser').get_text()
raw[:20]

#with open("text.txt", "w") as file:
#    file.write(raw)
```

```
Out[2]: "\n\nAll's Well That E"
```

1.2 Preprocess the Text

1.2.1 Split the text into tokens

```
In [3]: # word tokens
word_tokens = word_tokenize(raw)
sentence_tokens = sent_tokenize(raw)
```

```
In [4]: len(word_tokens)
```

```
Out[4]: 29364
```

```
In [5]: len(sentence_tokens)
```

```
Out[5]: 1442
```

1.2.2 Use NLTK Text to have only valid tokens

```
In [6]: text = nltk.Text(word_tokens)
```

```
In [7]: # (a) Find the frequency the word "BERTRAM" (all in caps, ignore font types such as bold, italics, etc.) that appears at the start of
text.count('BERTRAM')
```

```
Out[7]: 127
```

```
In [8]: # (b) Is the ratio of the frequency of BERTRAN Bertram
text.count('BERTRAM') / text.count('Bertram')
```

```
Out[8]: 18.142857142857142
```

```

In [9]: # (c) What is the average sentence length of this document?
total_len = 0
total_sent = 0
for sent in sentence_tokens:
    #print(sent)
    total_len += len(sent)
    total_sent += 1

print("total_len = {0}, total_sent = {1} , Avg Sent Len = {2}".format(total_len, total_sent, total_len / total_sent))

total_len = 129890, total_sent = 1442 , Avg Sent Len = 90.07628294036061

In [10]: # (d) What is the vocabulary (unique set of words) size? Ignore alpha- numeric and numeric terms, if available
vocab = [v.lower() for v in text.vocab() if v.isalpha()]
vocab = set(vocab)
print(len(vocab))
#print(vocab)

3299

In [11]: # This method doesn't work as expected

# (e) Find the total count of the words that end with the exclamation mark(!)
#TokenSearcher(raw).findall(r'<^[a-z][A-Z]+!>')
# print(len(TokenSearcher(raw).findall(r'[A-Z][a-z]\w*!')))

In [12]: # (e) Find the total count of the words that end with the exclamation mark(!)
regex_tokenizer = RegexpTokenizer('\w+')
regex_tokens = regex_tokenizer.tokenize(raw)
print(len(regex_tokens), regex_tokens)

131 ['tis!', 'living!', 'shape!', 'birthright!', 'head!', 'him!', 'you!', 'all!', 'queen!', 'monarch!', 'up!', 't!', 't!', 't!', 'well!', 'to']

```

2 Exercise 2 of Week 1 Assignment - Construct the binary incidence matrix

Construct the binary incidence matrix using the features extracted from the corpus. The corpus (271 text documents) is available at https://github.com/Ramaseshanr/anlp/blob/master/corpus/phy_corpus.txt. It contains contains questions from Kinematics class of physics problems sourced from the Internet - In this assignment, you need to develop a python program that uses the knowledge related to Kinematics and build a table similar to the one shown below for all the documents in the corpus. - The program should be able to read each problem, capture the known values (such as speed=10m/s, time=5s) and fill the respective cells in the table. For example, if you find 10 m/s for document 1, fill the speed with value row for D1 as 1. - Please note that problems may or may not contain all nine terms listed. - The corpus may contain duplicate entries - You may use any NLTK or any equivalent APIs for this assignment

	Terms	D1	D2	...	D271
Speed with value		1	0	...	0
Distance with value		0	0	...	1
Acceleration with value		0	0	...	0
Time with value		0	1	...	0

```

In [13]: import urllib.request
from bs4 import BeautifulSoup
import nltk
from nltk.tokenize import RegexpTokenizer # to get Kinematics
from nltk.tokenize import sent_tokenize # to find number of sentences in this corpus
from nltk.tokenize import word_tokenize
import pandas as pd # for table creation

```

2.1 Download the corpus

```

In [14]: url = 'https://raw.githubusercontent.com/Ramaseshanr/anlp/master/corpus/phy_corpus.txt'

# Download the corpus as text
html = urllib.request.urlopen(url).read().decode('utf-8')
soup = BeautifulSoup(html)
raw = BeautifulSoup(html, 'html.parser').get_text()
raw[:20]

# WRITE IT TO THE FILE if required
#with open("text.txt", "w", encoding="utf-8") as file:
#    file.write(raw)

```

Out[14]: 'An airplane accelera'

2.2 Preprocess the corpus

- Each Paragraph is considered as a separate document

```
In [15]: # Extract paragraphs
# Since text has indices in front of the lines, we can't use sent_tokenize()
# So splitting teh corpus into paragraphs and then I will extract kinematics from each paragraph
paragraphs = [para for para in raw.split('\n') if para]

In [16]: # Ensure we have right number of paragraphs that is found in the corpus
len(paragraphs)

Out[16]: 271
```

2.3 Find the various kind of units found in the given Corpus

```
In [17]: # Create RegEx Tokenizer which extracts the required kinematics
# Looking for 'x.xx letter' / 'x.xx word' / 'x.xx word\word'
#regex_tokenizer = RegexpTokenizer('|d+|.|d* |w |d+|.|d* [a-z]* |d+|.|d* [a-z]*/[a-z][0-9]* |w|.|d+ [a-z]*')
#pattern1 = '|d+|.|d+|s[a-zA-Z]*|/[a-zA-Z]*[d]*' # x{anything}xx X/X
#pattern1 = '|d+|. [da-zA-Z]+|s[a-zA-Z]*|/[a-zA-Z]*[d]*' # x{anything}xx X/X
pattern1 = '|d+|. [da-zA-Z]+|s[a-zA-Z]*|/[a-zA-Z\d\^~\-\]{1,}' # x{anything}xx X/X
pattern2 = '|d+|s[a-zA-Z]*|/[a-zA-Z]*[d]*' # x X/X
pattern3 = '|d+|.d+|s[a-zA-Z]+' # x{anything}xx X
pattern4 = '|d+|s[a-zA-Z]+' # x XX
pattern5 = '|.|.d+|s[a-zA-Z]+' # .x XX
regex_tokenizer = RegexpTokenizer(pattern1 + '|' + pattern2 + '|' + pattern3 + '|' + pattern4 + '|' + pattern5)
para_index = 0
total_tokens = 0
para_tokens = []
final_tokens = []
for paragraph in paragraphs:
    regex_tokens = regex_tokenizer.tokenize(paragraph)
    para_tokens.insert(para_index, [])
    if len(regex_tokens):
        final_tokens.extend(regex_tokens)
        para_tokens[para_index].extend(regex_tokens)
        print(para_index+1, regex_tokens)
        total_tokens += len(regex_tokens)
    para_index = para_index + 1
print('Total tokens found: ', total_tokens)
#print(final_tokens)
```

```
1 ['3.20 m/s2', '32.8 s']
2 ['5.21 seconds', '110 m']
3 ['2.60 seconds']
4 ['18.5 m/s', '46.1 m/s', '2.47 seconds']
5 ['1.40 meters', '1.67 m/s2']
6 ['444 m/s', '1.83 seconds']
7 ['7.10 m/s', '35.4 m']
8 ['3 m/s2', '65 m/s']
9 ['22.4 m/s', '2.55 s']
10 ['2.62 m']
11 ['1.29 m']
12 ['521 m/s', '0.840 m']
13 ['6.25 s']
14 ['370 m above']
15 ['367 m/s', '0.0621 m']
16 ['3.41 s']
17 ['290 m in', '3.90 m/s2']
18 ['88.3 m/s', '1365 m to']
19 ['112 m/s', '398 m']
20 ['1 m/s', '2.23 mi/hr', '91.5 m']
21 ['0.5 km', '1 hour later']
22 ['12 m/sec', '36 seconds']
23 ['2 m/s', '12 s he/she']
24 ['50 km traveling', '10 km/hr']
25 ['12 m/s', '3.00 minutes']
26 ['25 min at', '12 m/s']
27 ['3250 m/s', '10 m/s2', '215 km']
28 ['0.6 m/s2', '55 mi/h', '60 mi/h']
29 ['23.7 km/h', '0.92 m/s2', '3.6 s']
30 ['30 degree hill', '3.30 m/s2', '110 m']
31 ['48 m/s', '12 m/s', '5 s']
32 ['24 m/s', '315 m']
33 ['50 km/hr', '90 km/hr', '15 seconds']
34 ['9000 meters in', '12.12 seconds']
```

35 ['528 meters in', '4 seconds']
 36 ['3 minutes running', '6 m/s']
 37 ['450 km at', '120 m/s']
 38 ['1000 m in', '20 minutes']
 39 ['15 seconds', '0.8 m/sec', '7 m/sec']
 40 ['1.0 km/s', '1.8 km/s', '0.03 seconds']
 41 ['1700 m/s', '25 seconds']
 42 ['12 m/sec', '36 seconds']
 43 ['60 m/s', '8.0 sec']
 44 ['6.00 m/s²', '4.10 seconds']
 45 ['6 meter from']
 46 ['189 m/s', '20 m off']
 47 ['0.8 m/s', '10.3 seconds']
 48 ['3 seconds to']
 49 ['1 second it', '2 seconds']
 50 ['4 seconds', '3 seconds and']
 51 ['7 m/s', '20 seconds']
 52 ['115 meters high']
 53 ['110 m/s']
 55 ['29 m/s', '10 m/s']
 56 ['64 m in', '4 seconds']
 57 ['0)above the', '16 m/s', '10 m above']
 58 ['0.80 second']
 59 ['12.0 m/s', '70.0 m']
 60 ['10 m/s²', '4 seconds']
 61 ['6 m/s', '2 m/s²', '7 m from']
 62 ['16 meters from']
 63 ['30 m/s']
 64 ['1 second', '18 m/s', '3 seconds later']
 65 ['8 m/s²', '10 seconds', '10 seconds', '10 seconds']
 66 ['20 km/h', '8 m/s²', '10 seconds', '10 seconds', '10 seconds']
 67 ['0 to 72', '11.5 seconds', '72 km/h']
 68 ['40 m/s', '100 m']
 69 ['5 m/s', '20 m/s', '10 seconds', '10 seconds']
 70 ['350 km/h', '600 meters', '600 meters']
 71 ['360 km/h', '10 m/s²']
 72 ['8 seconds after', '340 m/s']
 73 ['10 m']
 74 ['10 m/s', '3.0 seconds']
 75 ['239 m high', '3.7 m/s²']
 76 ['1.40 m', '1.67 m/s²']
 77 ['2.6 seconds']
 78 ['1.29 m']
 79 ['420 m above']
 80 ['2 m/s']
 81 ['6 seconds of']
 82 ['9.0 seconds']
 83 ['2008 the Peachtree', '215 meters']
 84 ['15 seconds to', '2nd second', '5th second']
 85 ['3,000 meters']
 86 ['132 m high', '3.0 seconds', '3.0 seconds', '1.63 m/s²']
 87 ['0.45 meters']
 88 ['6 seconds of']
 89 ['9.0 seconds']
 90 ['2008 the Peachtree', '215 meters']
 91 ['1.85 seconds', '1.10 seconds']
 92 ['3.34 seconds']
 93 ['300.5 m']
 94 ['98.5 m/s']
 95 ['5.65 seconds']
 96 ['4.0 s', '4.0 s']
 97 ['1 instead threw', '8.0 m/s', '4.0 s', '8.0 m/s']
 98 ['1 instead threw', '8.0 m/s', '4.0 s', '8.0 m/s']
 99 ['8.0 m/s', '3.0 m/s']
 100 ['20.0 m', '10.0 m/s', '0 m/s']
 101 ['2.5 seconds']
 102 ['1.75 seconds']
 103 ['2.857 seconds']
 104 ['9.2 meters']
 105 ['1000 m']
 106 ['1 meter tall']
 107 ['9.2 meters']
 108 ['1000 m']
 109 ['1 meter tall']
 110 ['9.5 seconds']
 111 ['4 seconds to']

112 ['0.6 seconds']
 113 ['11 seconds later']
 114 ['9.5 seconds']
 115 ['0.6 seconds']
 116 ['11 seconds later']
 117 ['4 seconds to']
 118 ['18 meters']
 119 ['9000 meters when']
 120 ['1250 feet high', '1 foot', '0.305 meters']
 121 ['0.5 seconds', '2.00 m/s/s', '2.20 m/s/s']
 122 ['4.0 meters/second', '1 second', '1.5 m/s/s']
 123 ['48 m/s', '12 m/s', '5 s']
 124 ['24 m/s', '315 m']
 125 ['17.12 U', '24.55 Columbia', '25.71 U', '27.37 Wichita', '30.34 U', '2000 m long']
 126 ['45 miles straight', '75 mph', '85 mph', '20 minutes she']
 127 ['10⁶ m/s', '10¹⁴ m/s²', '10⁶ ms⁻¹']
 128 ['2.0 m/s', '2 until it', '20 m/s', '20 s at', '5.0 s']
 129 ['95 km/h', '130 km then', '65 km/h', '3 hours and', '20 min']
 130 ['89.5 km/h', '22 min rest', '77.8 km/h']
 131 ['0 to t', '4.21 min', '4.21 min', '8.42 min', '1.91 m/s', '1.00 min', '5.21 min']
 132 ['0.684 km', '0.486 km', '3.56 km', '61.7 degrees', '1.124 h']
 133 ['40 miles at', '80 mi/h', '40 mi/h']
 134 ['100-mile race', '50 mi/hr', '100 mi/hr']
 135 ['1.34 m/s', '6.44 km', '2.68 m/s', '0.447 m/s']
 136 ['6 mi/h', '5 mi', '5 mi in', '10 miles is', '10.8 mi/hr']
 137 ['1.50 km', '1.10 min']
 138 ['3250 m/s', '10 m/s²', '215 km']
 139 ['60 km/h', '2.0 s', '30 m and', '15 m wide', '6.4 m/s', '60 km/h', '70 km/h', '3.0 s']
 140 ['0.6 m/s²', '55 mi/h', '60 mi/h']
 141 ['23.7 km/h', '0.92 m/s', '2 for 3', '.6 s']
 142 ['112 m/s', '20.0 s']
 143 ['30 degree hill', '3.30 m/s', '110 m']
 144 ['0.30 km', '5.0 m/s', '33.0 m/s']
 145 ['24 m/s', '315 m']
 146 ['50 km/hr', '90 km/hr', '15 seconds']
 147 ['5.0 m/s', '4.0 seconds', '6.0 m/s']
 148 ['40 km/h', '13 m ahead', '8.0 m/s/s', '0.25 s']
 149 ['70 km/hr', '120 m']
 150 ['20 m/s', '3.5 s']
 151 ['20 seconds its', '108 kmh']
 152 ['20 m/s', '3.5 s']
 153 ['360 km/h', '24 seconds']
 154 ['1500 kilograms reaches', '15 metersvper second', '5 seconds after']
 155 ['0 to 60', '15 seconds']
 156 ['0 km/h', '100 km/h', '5 seconds', '5 m/s']
 157 ['150 km/h', '50 m from']
 158 ['3.20 m/s²', '32.8 s']
 159 ['18.5 m/s', '46.1 m/s', '2.47 seconds']
 160 ['3.30 m/s²', '88.0 m/s']
 161 ['8-km trip', '22.6 m/s', '16.8 m/s']
 162 ['25.0 m/s', '1.0 m/s²']
 163 ['25.0 m/s', '2.0 m/s²', '5 seconds', '10 more seconds', '15 seconds']
 164 ['0.5 mile/minute', '10 minutes', '.25 mile', '2 for 2', '12 minutes of']
 165 ['100 km at', '133 km/h']
 166 ['8 m/s', '0 if it', '3.0 seconds']
 167 ['9.58 s', '100-m dash']
 168 ['80 days', '40,075 km']
 169 ['3.20 m/s²', '32.8 s']
 170 ['5.21 seconds', '110 m']
 171 ['420 m above']
 172 ['2.6 seconds']
 173 ['120 cm east', '15 seconds', '120 cm west', '24 more seconds']
 174 ['380 km in', '60 km/hr']
 175 ['521 m/s', '0.840 m']
 176 ['367 m/s', '0.0621 m']
 177 ['88.3 m/s', '1365 m to']
 178 ['3.41 s']
 179 ['1.29 m']
 180 ['1.40 m', '1.67 m/s²']
 181 ['0.926 m/s²', '15 seconds with', '15 m/s']
 182 ['2.5 meters']
 183 ['2.8 m']
 184 ['23.7 km/h', '0.92 m/s²', '3.6 s']
 185 ['4.30 m/s', '3.0 m/s²', '5.0 s']
 186 ['5.0 s', '1.5 m/s²']
 187 ['15.0 m/s', '2.0 m/s²', '10.0 m/s']

188 ['2.3 m/s', '55 m']
 189 ['120 km/hr', '240 m']
 190 ['32 m/s', '3.0 m/s', '2 for 9']
 191 ['7.0 m/s', '0.80 m/s²', '245 m', '125 m', '67 m']
 192 ['15 m/s', '2 m/s²', '2.5 seconds']
 193 ['5 seconds with', '1.5 m/s²']
 194 ['4.30 m/s', '3 m/s²', '5 seconds']
 195 ['23.7 km/hr', '0.92 m/s²', '3.6 sec']
 196 ['90 km/hr', '36 km/hr', '5 s']
 197 ['90 km/hr', '10 seconds to', '36 km/hr', '15 s']
 198 ['25 m/s', '5 m/s', '4 s']
 199 ['2 m/s', '18 m/s', '8 s']
 200 ['3 m/s', '5 m/s', '6 s', '30 m/s']
 201 ['500 m due', '300 s and', '400 m in', '320 s in', '500 m', '400 m and']
 202 ['2 towns 60', '2 hours']
 203 ['2 meters/second', '4 minutes']
 204 ['6 m/s', '14 m/s²', '7e6 m/s']
 205 ['0.210 s', '1.35 meters']
 206 ['6.5 seconds']
 207 ['3 seconds before', '3 second drop']
 208 ['7.0 m', '2.00 m/s']
 209 ['2.0 m/s', '9.8 m/s', '2.5 m']
 210 ['4.6 m/s']
 211 ['460 m/s']
 212 ['7 seconds after', '50 m/s']
 213 ['24 m above']
 214 ['1.50 m', '19.6 m/s']
 215 ['2 m/s', '50 m above']
 216 ['5.00 m/s', '105 m above']
 217 ['9.8 m/s']
 218 ['30 m/s', '4.0 seconds', '10 m/s']
 219 ['12 m/s', '40 degrees above']
 220 ['12 m/s', '30.0 m']
 221 ['54 m high', '130 m from']
 222 ['22.2 m/s', '36 m from']
 223 ['6.5 m', '4.0 m/s']
 224 ['215 km/h', '155 km/h', '78.0 m']
 225 ['25 m/s']
 227 ['2.5 m/s', '0.75 m/s', '4.0 s']
 228 ['30.0 m']
 229 ['10 m', '15 degrees']
 230 ['38.9 m', '3.05 m', '20.4 m/s']
 231 ['30 m/s', '2 near Earth']
 232 ['6.8 m/s', '2 m away']
 233 ['5.0 kg', '10.1 m', '8.6 m/s']
 234 ['48 m above', '24 m/s']
 235 ['125 m above', '65.0 m/s']
 237 ['3.05 m', '2 m above', '10 m from', '45 degree angle']
 238 ['8.0 m', '9.0 m']
 239 ['15.0 m/s', '30 degrees above', '2.0 seconds']
 240 ['25 m/s']
 241 ['26 m away', '5 m/s']
 242 ['20.2 m', '10 degrees above', '3.0 seconds']
 243 ['2.00 m', '.55 m', '32 degrees']
 244 ['301.5 m', '25 degree to']
 245 ['3.00 kg', '176.4 m', '12.0 N']
 246 ['60 degrees with', '5 seconds later']
 247 ['30 m away', '5 m above', '3 seconds after']
 248 ['11 meters']
 249 ['32.5 m/s']
 250 ['10 m']
 252 ['41.8 m', '42.7 m/s', '33.0 degrees']
 254 ['120 m/s', '55 m/s']
 255 ['35 m down', '3.06 m/s', '45 m high']
 256 ['00-kg rock', '5.00 m/s']
 257 ['58 m and', '3.41 s']
 258 ['20 meters and', '10 seconds to']
 259 ['5 m/s', '10 m above']
 260 ['4.9 m/s', '2 s', '9.8 m/s²']
 261 ['10 m/s²']
 262 ['8 s in', '50 m']
 263 ['75 m tall', '150 m away']
 264 ['2.64 m/s', '1.48 s']
 265 ['54 m high', '130 m from']
 266 ['10 m/s']
 267 ['20 m/s²', '1.0 min']

```

268 ['150000000 km', '3.0E8 m/s']
269 ['2.15 is', '15 m/s', '18 m']
270 ['0.0 m/s', '4.0 m/s', '4.0 s']
271 ['73.5 m/s', '2.2 s']
Total tokens found: 600

```

```

In [18]: def is_number_repl_dot_isdigit(s):
        """
        replaces '.' with '' in s and
        returns the result of isdigit(s)
        https://stackoverflow.com/questions/354038/how-do-i-check-if-a-string-is-a-number-float
        """
        return s.replace('.', '', 1).isdigit()

```

```

In [19]: def get_tokens(final_tokens):
        """
        Accepts list having strings
        splits into tokens separated by spaces
        and returns it as a set
        """
        ll = []
        [ll.extend(token.split(' ')) for token in final_tokens]
        return set(ll)

```

```

In [20]: final_tokens = get_tokens(final_tokens)
        final_tokens = [s for s in final_tokens if not is_number_repl_dot_isdigit(s)]
        sorted(final_tokens)

```

```

Out[20]: ['0)above',
          '00-kg',
          '100-m',
          '100-mile',
          '10^14',
          '10^6',
          '2nd',
          '3,000',
          '3.0E8',
          '40,075',
          '5th',
          '7e6',
          '8-km',
          'Columbia',
          'Earth',
          'N',
          'Peachtree',
          'U',
          'Wichita',
          'above',
          'after',
          'ahead',
          'and',
          'angle',
          'at',
          'away',
          'before',
          'cm',
          'dash',
          'days',
          'degree',
          'degrees',
          'down',
          'drop',
          'due',
          'east',
          'feet',
          'foot',
          'for',
          'from',
          'h',
          'he/she',
          'high',
          'hill',
          'hour',
          'hours',
          'if',
          'in',

```

```

'instead',
'is',
'it',
'its',
'kg',
'kilograms',
'km',
'km/h',
'km/hr',
'km/s',
'kmh',
'later',
'long',
'm',
'm/s',
'm/s/s',
'm/s^2',
'm/s^2',
'm/sec',
'meter',
'meters',
'meters/second',
'metersvper',
'mi',
'mi/h',
'mi/hr',
'mile',
'mile/minute',
'miles',
'min',
'minutes',
'more',
'mph',
'ms^-1',
'near',
'of',
'off',
'race',
'reaches',
'rest',
'rock',
'running',
's',
'sec',
'second',
'seconds',
'she',
'straight',
't',
'tall',
'the',
'then',
'threw',
'to',
'towns',
'traveling',
'trip',
'until',
'west',
'when',
'wide',
'with']

```

2.4 Create the untills table based on above finding

```

In [21]: speed_with_val = ['minutes', 'min', 'mi/hr', 'mi/h', 'mile/minute', 'metersvper', 'meters/second', 'm/sec', 'm/s', 'ms^-1', 'mph']
distance_with_val = ['miles', 'mile', 'meters', 'meter', 'm', 'feet', '-km', '-mile', '-m', 'km']
acceleration_with_val = ['km/h', 'km/hr', 'km/s', 'kmh', 'm/s^2', 'm/s^2', 'm/s/s']
time_with_val = ['hours', 'hour', 'second', 'seconds', 'sec', 's', 'times', 'h']

```

2.5 Create the table and fill it

```

In [22]: column_header = ['Speed with value', 'Distance with value', 'Acceleration with value', 'Time with value']
df = pd.DataFrame(columns=column_header)

```

```

In [23]: # loop through each paragraph tokens

```



```

para_index = 0
for tokens in para_tokens:
    tt = get_tokens(tokens)
    #print(tt)
    # loop through each token
    c1 = 0
    c2 = 0
    c3 = 0
    c4 = 0
    if any(s in tt for s in speed_with_val):
        c1 = 1
    if any(s in tt for s in distance_with_val):
        c2 = 1
    if any(s in tt for s in acceleration_with_val):
        c3 = 1
    if any(s in tt for s in time_with_val):
        c4 = 1
    #print(para_index+1, c1, c2, c3, c4)
    df = df.append(pd.Series([c1, c2, c3, c4], index=df.columns), ignore_index=True)
para_index += 1

```

In [24]: df

```

Out[24]:   Speed with value  Distance with value  Acceleration with value  \
0           0           0           1
1           0           1           0
2           0           0           0
3           1           0           0
4           0           1           1
..          ...          ...          ...
266          1           0           1
267          1           1           0
268          1           1           0
269          1           0           0
270          1           0           0

      Time with value
0           1
1           1
2           1
3           1
4           0
..          ...
266          0
267          0
268          0
269          1
270          1

[271 rows x 4 columns]

```

In [25]: df.transpose()

```

Out[25]:   0  1  2  3  4  5  6  7  8  9  ...  261 262  \
Speed with value  0  0  0  1  0  1  1  1  1  0  ...  0  0
Distance with value  0  1  0  0  1  0  1  0  0  1  ...  1  1
Acceleration with value  1  0  0  0  1  0  0  1  0  0  ...  0  0
Time with value    1  1  1  1  0  1  0  0  1  0  ...  1  0

      263 264 265 266 267 268 269 270
Speed with value    1  0  1  1  1  1  1  1
Distance with value    0  1  0  0  1  1  0  0
Acceleration with value  0  0  0  1  0  0  0  0
Time with value      1  0  0  0  0  0  1  1

[4 rows x 271 columns]

```

3 Exercise 3 of Week 1 Assignment

Write a program to find out whether Mandelbrot's approximation really provides a better fit than Zipf's empirical law. Use the same corpus for Zipf and Mandelbrot approximation

- Zipf Law
- <https://www.cs.swarthmore.edu/~richardw/classes/cs65/f18/lab01.html> for loglog
- https://www.researchgate.net/publication/221200132_Exploring_Regularity_in_Source_Code_Software_Science_and_Zipf's_Law

```

In [26]: from operator import itemgetter
import nltk
# from nltk.probability import FreqDist
from nltk.corpus import stopwords

import matplotlib.pyplot as plt
import math

In [27]: stop_words = set(stopwords.words('english'))

words = nltk.Text(nltk.corpus.gutenberg.words('austen-emma.txt'))
# words = nltk.Text(nltk.corpus.gutenberg.words('carroll-alice.txt'))

# convert to lower case and consider only words
words = [word.lower() for word in words] #if word.isalpha()
# remove stop words
# words = [word for word in words if word not in stop_words]

In [28]: # =====
# fDist = FreqDist(words[:2000])
#
# print(fDist.B())
# print(fDist.Nr(1153))
#
# for k,v in fDist.items():
#     print(k, v)
#
# fDist.plot()
# =====

In [29]: # build word frequency as dictionary
frequency = {}
for word in words:
    count = frequency.get(word, 0)
    frequency[word] = count + 1

# sort and create a list
freq_list = sorted(frequency.items(), key=itemgetter(1), reverse=True)
#print(type(freq_list), freq_list)

In [30]: # plot loglog graph of rank versus frequency
ranks = range(1, len(freq_list)+1) # x-axis: ranks
freqs = [freq for (word, freq) in freq_list] # y-axis: frequencies

print('50th Elem Freq: ', freqs[50])
print('150th Elem Freq: ', freqs[150])
print('150th*3 ==> ', freqs[150] * 3)

plt.loglog(ranks, freqs, label='austen-emma.txt')
plt.xlabel('log(rank)')
plt.ylabel('log(freq)')
plt.legend(loc='lower left')

```

```

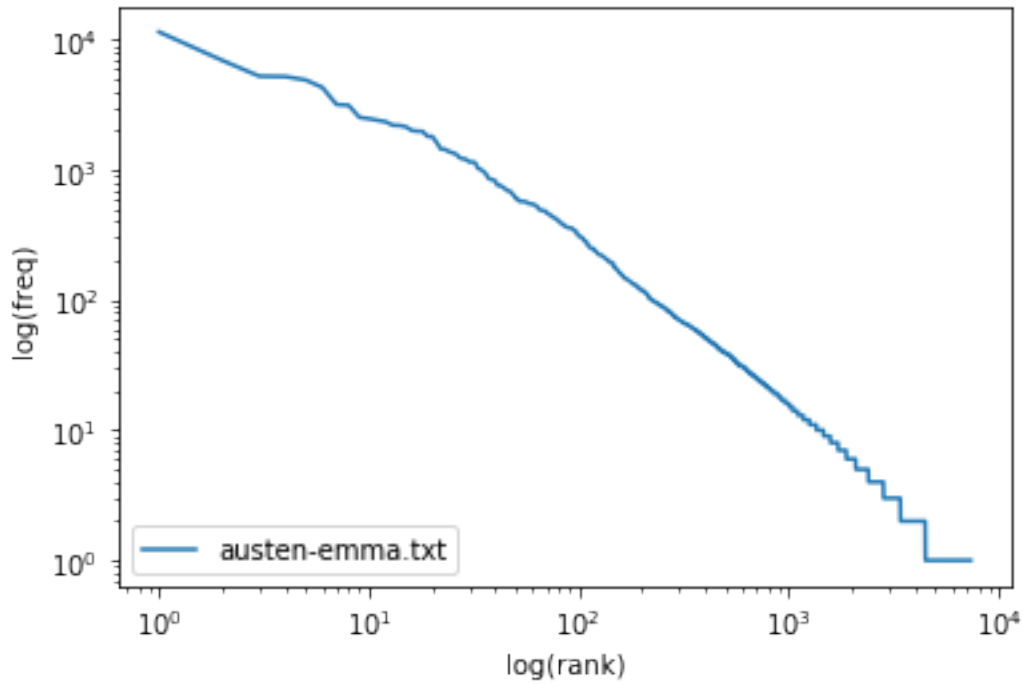
50th Elem Freq:  599
150th Elem Freq:  173
150th*3 ==>  519

```

```

Out[30]: <matplotlib.legend.Legend at 0x2d27382fcf8>

```




```
In [31]: # plot zip law's expected value
# https://www.cs.swarthmore.edu/~richardw/classes/cs65/f18/lab01.html
def H_approx(n):
    """
    Returns an approximate value of n-th harmonic number
    http://en.wikipedia.org/wiki/Harmonic-number
    """
    # Euler-Mascheroni constant
    gamma = 0.57721566490153286060651209008240243104215933593992
    return gamma + math.log(n) + 0.5/n - 1./(12*n**2) + 1./(120*n**4)

T = len(words)
print('T: ', T)
vocab = set([word for word, freq in freq_list])
n = len(vocab)
print('n: ', n)
k = T/H_approx(n)
print('k: ', k)
expected_freq = [k/r for r in ranks]
plt.loglog(ranks, expected_freq, label="zipf's law")
plt.legend(loc='lower left')

plt.show()
```

```
T: 192427
n: 7344
k: 20300.51373217429
```



4 Exercise 4 of Week 1 Assignment

```
In [32]: import nltk
         from nltk.corpus import stopwords
         import pandas as pd

         stop_words = stopwords.words('english')

In [33]: def getHeapLawValues(corpus_name):
         """
         Accepts nltk.corpus.gutenberg corpus name
         Calculated vocabulary (M) and total words (T)
         Returns M, T, M/T as a list
         """
         words = nltk.Text(nltk.corpus.gutenberg.words(corpus_name))
         # normalize the words
         words = [w.lower() for w in words if w.isalpha()]
         # remove stop words
         words = [w for w in words if w not in stop_words]

         M = len(set(words))
         T = len(words)

         print(corpus_name, ' M: ', M, ' T: ', T, ' Ratio: ', M/T)
         return [M, T, M/T]

In [34]: column_headers = ['M', 'T', 'Ratio']
         df = pd.DataFrame(columns=column_headers)

         for corpus in nltk.corpus.gutenberg.fileids():
             df.loc[corpus] = getHeapLawValues(corpus)

         print(df)
         #df.plot()
         df.Ratio.plot()

austen-emma.txt M: 6948 T: 73149 Ratio: 0.09498421031046221
austen-persuasion.txt M: 5606 T: 38337 Ratio: 0.1462294910921564
austen-sense.txt M: 6148 T: 53986 Ratio: 0.11388137665320638
bible-kjv.txt M: 12443 T: 374945 Ratio: 0.03318620064276094
blake-poems.txt M: 1400 T: 3805 Ratio: 0.3679369250985545
bryant-stories.txt M: 3688 T: 21718 Ratio: 0.16981305829266047
burgess-busterbrown.txt M: 1382 T: 7582 Ratio: 0.18227380638353996
carroll-alice.txt M: 2423 T: 12240 Ratio: 0.19795751633986927
chesterton-ball.txt M: 8009 T: 39715 Ratio: 0.20166184061437745
chesterton-brown.txt M: 7589 T: 35348 Ratio: 0.21469390064501528
```

```

chesterton-thursday.txt M: 6159 T: 28328 Ratio: 0.21741739621575826
edgeworth-parents.txt M: 8166 T: 78148 Ratio: 0.1044940369555203
melville-moby_dick.txt M: 16802 T: 110459 Ratio: 0.15211073792085752
milton-paradise.txt M: 8849 T: 45568 Ratio: 0.19419329353932585
shakespeare-caesar.txt M: 2911 T: 11056 Ratio: 0.2632959479015919
shakespeare-hamlet.txt M: 4590 T: 15898 Ratio: 0.288715561705875
shakespeare-macbeth.txt M: 3340 T: 10078 Ratio: 0.3314149632863663
whitman-leaves.txt M: 12189 T: 65080 Ratio: 0.18729256299938538

```

	M	T	Ratio
austen-emma.txt	6948.0	73149.0	0.094984
austen-persuasion.txt	5606.0	38337.0	0.146229
austen-sense.txt	6148.0	53986.0	0.113881
bible-kjv.txt	12443.0	374945.0	0.033186
blake-poems.txt	1400.0	3805.0	0.367937
bryant-stories.txt	3688.0	21718.0	0.169813
burgess-busterbrown.txt	1382.0	7582.0	0.182274
carroll-alice.txt	2423.0	12240.0	0.197958
chesterton-ball.txt	8009.0	39715.0	0.201662
chesterton-brown.txt	7589.0	35348.0	0.214694
chesterton-thursday.txt	6159.0	28328.0	0.217417
edgeworth-parents.txt	8166.0	78148.0	0.104494
melville-moby_dick.txt	16802.0	110459.0	0.152111
milton-paradise.txt	8849.0	45568.0	0.194193
shakespeare-caesar.txt	2911.0	11056.0	0.263296
shakespeare-hamlet.txt	4590.0	15898.0	0.288716
shakespeare-macbeth.txt	3340.0	10078.0	0.331415
whitman-leaves.txt	12189.0	65080.0	0.187293

```

Out[34]: <matplotlib.axes._subplots.AxesSubplot at 0x2d2717a22b0>

```

